

02409 Multivariate Statistics

Lecture I, November 2 2025

Anders Stockmarr

anst@dtu.dk

(1-3) 60%

Clustering 4 groups

Course developers:

Anders Stockmarr

Anders Nymark Christensen



Factor 1 [41%]

Factor 3 [19%]



Agenda

- GLM in practice;
- Model selection;
- Outliers and influential observations

The General Linear Model - GLM

\boldsymbol{x} and $\boldsymbol{\Sigma}$ assumed known:

$$\boldsymbol{Y} \sim N(\boldsymbol{\mu}, \sigma^2 \boldsymbol{\Sigma}) \quad \boldsymbol{\mu} = \mathbf{x} \boldsymbol{\theta}$$

$$\begin{bmatrix} \mu_1 \\ \vdots \\ \mu_n \end{bmatrix} = \begin{bmatrix} x_{11} & \cdots & x_{1k} \\ \vdots & & \vdots \\ x_{n1} & \cdots & x_{nk} \end{bmatrix} \begin{bmatrix} \theta_1 \\ \vdots \\ \theta_k \end{bmatrix}$$

or

$$\begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix} = \begin{bmatrix} x_{11} & \cdots & x_{1k} \\ \vdots & & \vdots \\ x_{n1} & \cdots & x_{nk} \end{bmatrix} \begin{bmatrix} \theta_1 \\ \vdots \\ \theta_k \end{bmatrix} + \begin{bmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_n \end{bmatrix}$$

i.e.

$$\boldsymbol{Y} = \mathbf{x} \boldsymbol{\theta} + \boldsymbol{\varepsilon}$$

with

$$\boldsymbol{\varepsilon} \sim N(\mathbf{0}, \sigma^2 \boldsymbol{\Sigma}), \quad \mathbf{x}, \boldsymbol{\Sigma} \text{ known}$$

Estimation in GLM I

||| Theorem 2.3

If Y is independent and identical distributed (iid), Σ is the identity matrix!

Let x and θ be given as in the preceding section and let $Y \sim N_n(x\theta, \sigma^2\Sigma)$, where Σ is positive definite. Then the maximum likelihood estimator $\hat{\theta}$ for θ is given by $x\hat{\theta}$ being the projection (with respect to Σ) onto M , $\hat{\theta}$ is a solution to the so-called *normal equation(s)*

$$(x^T \Sigma^{-1} x) \hat{\theta} = x^T \Sigma^{-1} y.$$

If Σ is the identity matrix,
we simply have ordinary
least squares

If x has full rank k , then

$$\Sigma = I \Rightarrow \hat{\theta} = (x^T x)^{-1} x^T Y$$

$$\hat{\theta} = (x^T \Sigma^{-1} x)^{-1} x^T \Sigma^{-1} Y,$$

and being a linear combination of normally distributed variables $\hat{\theta}$ is also normally distributed with parameters

$$\begin{aligned} E(\hat{\theta}) &= \theta \\ D(\hat{\theta}) &= \sigma^2 (x^T \Sigma^{-1} x)^{-1}. \end{aligned}$$

It is especially noted that $\hat{\theta}$ is an unbiased estimate of θ .

Estimation in GLM II

||| Theorem 2.5

Let the situation be as above. The maximum likelihood estimator of σ^2 is

$$\tilde{\sigma}^2 = \frac{1}{n} \|Y - x\hat{\theta}\|^2 = \frac{1}{n} (Y - x\hat{\theta})^T \Sigma^{-1} (Y - x\hat{\theta}).$$

The unbiased estimator of σ^2 is

If Σ is the identity matrix, we simply have the equations from ordinary least squares

$$\begin{aligned}\hat{\sigma}^2 &= \frac{1}{n - \text{rk}(x)} \|Y - x\hat{\theta}\|^2 \\ &= \frac{1}{n - \text{rk}(x)} (Y - x\hat{\theta})^T \Sigma^{-1} (Y - x\hat{\theta})\end{aligned}$$

where $x\hat{\theta}$ is the maximum likelihood estimator of $E(Y)$. The following holds

$$\hat{\sigma}^2 \sim \sigma^2 \chi^2(n - \text{rk}(x)) / (n - \text{rk}(x))$$

and $\hat{\sigma}^2$ is independent of the maximum likelihood estimator of the expected value and is therefore independent of $\hat{\theta}$.

Example: Crime data

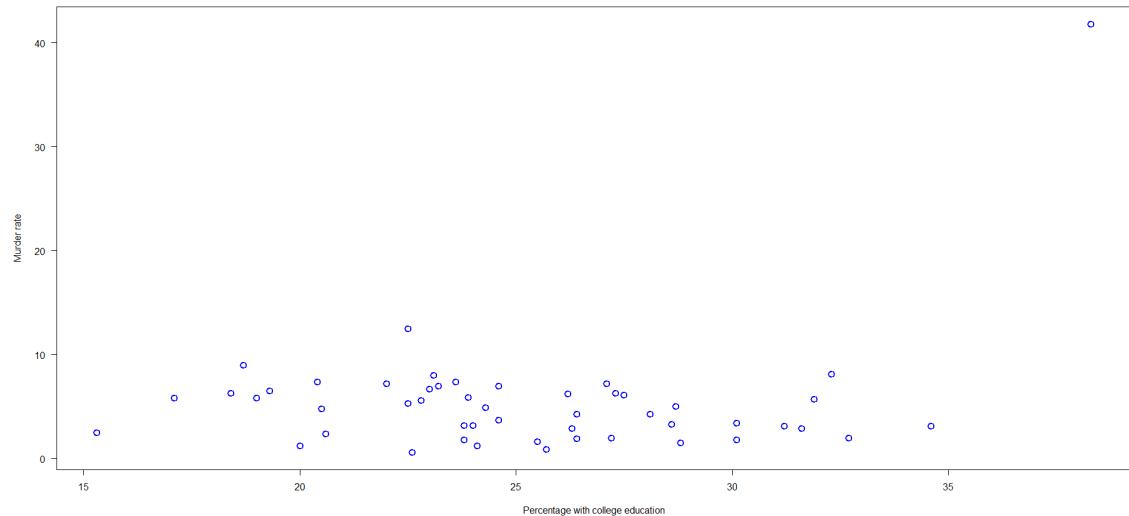
- Is higher education associated with higher murder rates?
- Crime data from 2003 for the US.
- **Murder rate:** The annual number of murders per 100,000 people in the population.
- **College:** Percentage of the adult residents who have a college education.
- Y is the murder rate; X is the explanatory variable college.

Example: Crime data

```
> crime <- read.delim("Data/us_statewide_crime.txt")
> crime<-crime[,c(1,3,6)]
> head(crime)

      State murder.rate college
1    Alabama      7.4     20.4
2     Alaska      4.3     28.1
3   Arizona      7.0     24.6
4  Arkansas      6.3     18.4
5 California      6.1     27.5
6 Colorado       3.1     34.6

> plot(crime$college, crime$murder.rate,
+       xlab="Percentage with college education",
+       ylab="Murder rate", las=1, cex = 1.5,
+       col = "blue", lwd = 2)
```



Example: Crime data

- We set up a linear model: With Y the murder rate,

$$Y_i = \alpha + \beta X_i + \varepsilon_i, \quad \varepsilon_i \sim iid. N(0, \sigma^2),$$

Or

$$Y = X\theta + \varepsilon,$$

$$X = \begin{bmatrix} 1 & college_1 \\ \vdots & \vdots \\ 1 & college_{51} \end{bmatrix}, \theta = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}$$

Why is this a good idea?

Taylor expansion of real functions

- Unknown relationship between crime and college:

$$Y = f(X) + \varepsilon$$

Power series expansion:

$$f(x) = f(x_0) + \sum_{n=1}^{\infty} a_n (x - x_0)^n, \quad a_n = \frac{f^{(n)}(x_0)}{n!}$$
$$f(x) = f(x_0) + f'(x_0)(x - x_0) + r_1(x - x_0)$$

$$= \underbrace{f(x_0) - f'(x_0)x_0}_{\alpha} + \underbrace{f'(x_0)}_{\beta} x + r_1(x - x_0)$$

So that

$$Y = \alpha + \beta X + r_1(x - x_0) + \varepsilon = \alpha + \beta X + \tilde{\varepsilon}$$

IF the remainder term $r_1(x - x_0)$ is negligible compared to the residual ε . **Note:** $r_1(x - x_0) = o(x - x_0)$.

Taylor expansion of real functions

"**IF** the remainder term $r_1(x - x_0)$ is negligible compared to the residual ε "

- What if that is not the case?
- **Taylor expansion of order 2:**

$$Y = f(X) + \varepsilon$$

Power series expansion (higher dimensions of X):

$$f(x) = f(x_0) + \langle f'(x_0), x - x_0 \rangle + \frac{1}{2} \|x - x_0\|_{f''(x_0)}^2 + r_2(x - x_0)$$

$$Y = \alpha + \sum_{i=1}^k \beta_i X_i + \sum_{r,s=1}^k \beta_{rs} X_r X_s + \tilde{\varepsilon}$$

IF the remainder term $r_2(x - x_0)$ is negligible compared to the residual ε . **Note:** $r_2(x - x_0) = o((x - x_0)^2)$.

Example: Crime data

- We set up a linear model:

$$Y_i = \alpha + \beta X_i + \varepsilon_i, \quad \varepsilon_i \sim iid. N(0, \sigma^2),$$

Or

$$Y = X\theta + \varepsilon,$$

$$X = \begin{bmatrix} 1 & college_1 \\ \vdots & \vdots \\ 1 & college_{51} \end{bmatrix}, \theta = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}$$

- 1st order Taylor approximation to the relationship between murder rate and college.

Example: Crime data

- Parameter estimates:

Note that

$$X^T X = \begin{bmatrix} n & co. \\ co. & co^2 \end{bmatrix}$$

Where a dot indicates summation. Thus:

$$(X^T X)^{-1} = \frac{1}{nco^2 - co.^2} \begin{bmatrix} co^2. & -co. \\ -co. & n \end{bmatrix} = \frac{1}{SSD_{co}} \begin{bmatrix} \overline{co^2} & -\overline{co} \\ -\overline{co} & 1 \end{bmatrix}$$

where a overbar indicates average, and where $SSD_{co} = \sum_i (co_i - \overline{co})^2$.

SSD_{co} = Sum of Squares of Deviations of college.

Example: Crime data

- Parameter estimates:

$$\begin{pmatrix} \hat{\alpha} \\ \hat{\beta} \end{pmatrix} = (X^T X)^{-1} X^T Y = \frac{1}{SSD_{co}} \begin{bmatrix} \overline{co^2} & -\overline{co} \\ -\overline{co} & 1 \end{bmatrix} \begin{bmatrix} cr. \\ < co, cr > \end{bmatrix}$$

= ...

$$= \left[\frac{cr. - \hat{\beta} \overline{co}}{SSD_{co}} \right] = \left[\frac{cr. - \hat{\beta} \overline{co}}{SPD_{co,cr}} \right]$$

with $SPD_{co,cr}$ = Sum of Products of Deviations of college and crime.

Example: Crime data

- Parameter estimates in **R**:

```
> x <- crime$college
> y <- crime$murder.rate
> (beta <- sum( (x-mean(x)) * (y-mean(y)) ) / sum( (x-mean(x))^2 ) )
[1] 0.33307

> (alpha <- mean(y) - beta*mean(x))
[1] -3.058072

> Fitted <- alpha + beta*x; Resid <- y - Fitted
> sigma2 <- sum(Resid^2) / (length(y)-2)
> sqrt(sigma2)
[1] 5.6146
```

- Thus:

$$\hat{\alpha} = -3.058072, \quad \hat{\beta} = 0.33307, \quad \hat{\sigma} = 5.6146$$

Example: Crime data

- Parameter estimates in R using the lm() function:

```
> reg1 <- lm(murder.rate ~ college, data = crime)
> summary(reg1)
```

Call:
lm(formula = murder.rate ~ college, data = crime)

Residuals:
Min 1Q Median 3Q Max
-5.833 -3.368 -1.403 1.981 32.101

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) -3.0581 4.3630 -0.701 0.4867
college 0.3331 0.1703 1.955 0.0563 .

Signif. codes: 0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1

Residual standard error: 5.615 on 49 degrees of freedom
Multiple R-squared: 0.07238, Adjusted R-squared: 0.05345
F-statistic: 3.823 on 1 and 49 DF, p-value: 0.05626

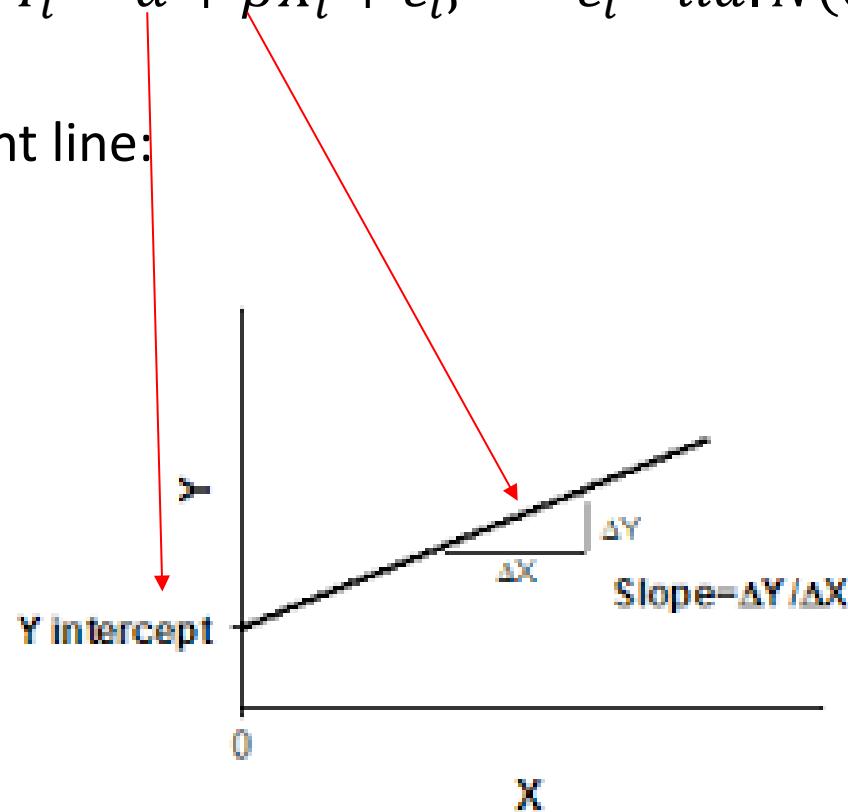
$$y = \hat{\alpha} + \hat{\beta}x$$

$$\hat{\sigma}$$

Example: Crime data

$$Y_i = \alpha + \beta X_i + \varepsilon_i, \quad \varepsilon_i \sim iid. N(0, \sigma^2)$$

The expression for a straight line:



The slope β is how much the murder rate increases when the college rate increase by one unit.
 $\hat{\beta} = 0.3331$.

Example: Crime data

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|----------|
| (Intercept) | -3.0581 | 4.3630 | -0.701 | 0.4867 |
| college | 0.3331 | 0.1703 | 1.955 | 0.0563 . |

- Parameter estimates handing in R:

```
> confint(reg1)
              2.5 %    97.5 %
(Intercept) -11.825900553 5.7097573
college      -0.009244278 0.6753843

> #Nice table
> tab <- cbind(coef(summary(reg1))[, 1:2], "Lower" = confint(reg1)[, 1],
+                 "Upper" = confint(reg1)[, 2])

> tab
      Estimate Std. Error       Lower       Upper
(Intercept) -3.058072  4.3630260 -11.825900553 5.7097573
college      0.333070  0.1703416 -0.009244278 0.6753843

> #Nice table with p-values
> data.frame(round(tab, 2),
+             "p-value" = format.pval(coef(summary(reg1))[, 4], digits = 3, eps = 1e-3))
      Estimate Std..Error Lower Upper p.value
(Intercept)    -3.06        4.36 -11.83  5.71  0.4867
college        0.33        0.17  -0.01  0.68  0.0563
```

Example: Crime data

- **Model check:**
 - Results can only be trusted if the model fits the data.

$$Y_i = \alpha + \beta X_i + \varepsilon_i, \quad \varepsilon_i \sim iid. N(0, \sigma^2)$$

- **Model assumptions:**
 1. Linear dependency of covariates;
 2. Independence of residuals;
 3. Normality of residuals;
 4. Homogeneity of residuals (same σ^2 for all observations).

Example: Crime data

1. Linear dependency of covariates;
 - scatter plots of residuals vs. numerical explanatory variables;
2. Independence of residuals;
 - Context; cannot be checked with numerical methods in simple models;
3. Normality of residuals;
 - Quantile-Quantile plots of residuals versus the normal distribution;
4. Homogeneity of residuals (same σ^2 for all observations).
 - scatter plots of residuals vs. fitted values.

Example: Crime data

Quantile-quantile plots:

- Two distributions P and Q are the same if their distribution functions are the same:

$$\forall x \in \mathbb{R}: P((-\infty; x]) = Q((-\infty; x])$$

- Similar: Two distributions P and Q are the same if their inverse distribution functions (**quantile functions**) are the same:

$$\begin{aligned}\forall p \in [0; 1]: q_{P,p} &= \inf_{x \in \mathbb{R}} \{x: P((-\infty; x]) \geq p\} \\ &= \inf_{x \in \mathbb{R}} \{x: Q((-\infty; x]) \geq p\} = q_{Q,p}\end{aligned}$$

Check of normality: Compare empirical quantiles with the quantiles of the corresponding normal distribution

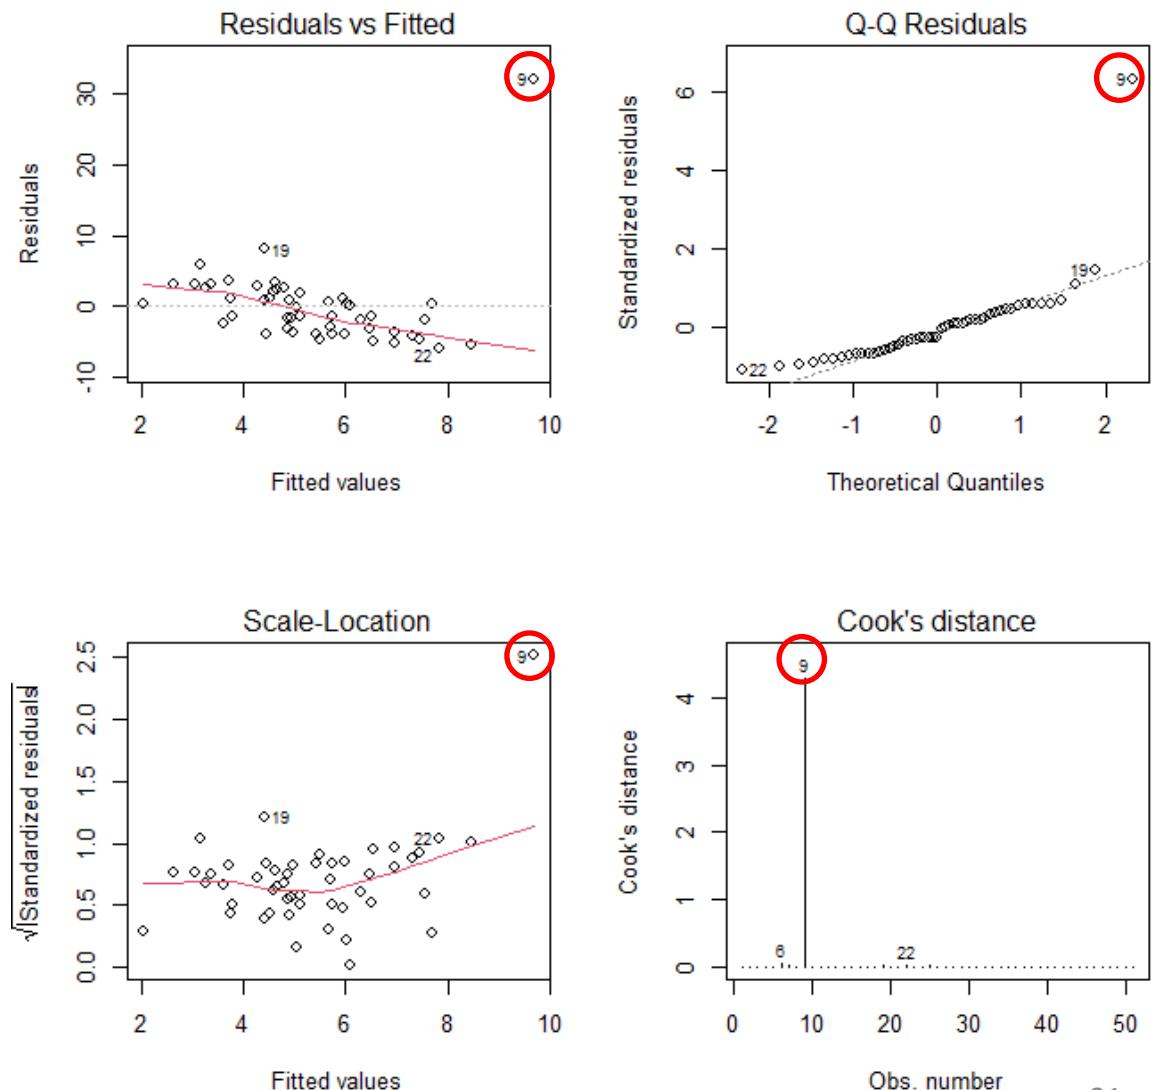
Example: Crime data

```
par(mfrow = c(2, 2))
plot(reg1, which = 1:4)
par(mfrow = c(1, 1))
```

There is an issue with observation #9!

```
> crime[9,]
      State murder.rate college
9 District of Columbia     41.8    38.3
```

Washington DC, the capital district of USA –
not a state.



Example: Crime data

Outliers:

Incorrect data; OR data generated from a data generating mechanism deviating from the rest of the model.

Must be checked, corrected, or removed.

We want to uncover the underlying data generating mechanism; observations not complying with that will

- a) disturb the interpretation and
- b) not contribute with information.

Example: Crime data

Observation #9 is an outlier:

Observed value: 41.8.

Predicted value:

$$\hat{\alpha} + \hat{\beta} * 38.3 = -3.058072 + 0.33307 * 38.3 = 9.69851$$

Estimated residual: $41.8 - 9.69851 = 32.10149$

$$\begin{aligned}\hat{P}\left(\max_{i=1,\dots,51} |\varepsilon_i| \geq 32.10149\right) &= 1 - \hat{P}\left(\max_{i=1,\dots,51} |\varepsilon_i| < 32.10149\right) \\ &= 1 - \hat{P}(|\varepsilon_i| < 32.10149)^{51} 1 - \left(1 - 2\Phi\left(\frac{-32.10149}{\hat{\sigma}}\right)\right)^{51} \\ &= 0.0000006\end{aligned}$$

Example: Crime data

- Observation #9 is also the observation with the biggest residual;
The observation that contributes most to the estimated variation.
- What if we estimated the variation disregarding the suspected outlier?

```
>crime50 <- crime[-9, ]  
>reg2 <- lm(murder.rate ~ college, data = crime50)
```

```
>(sigma.new2<-summary(reg2)$sigma^2)  
[1] 2.464061
```

```
>1-(1-2*(1-pnorm(reg1$res[9]/sigma.new)))^51
```

9

0

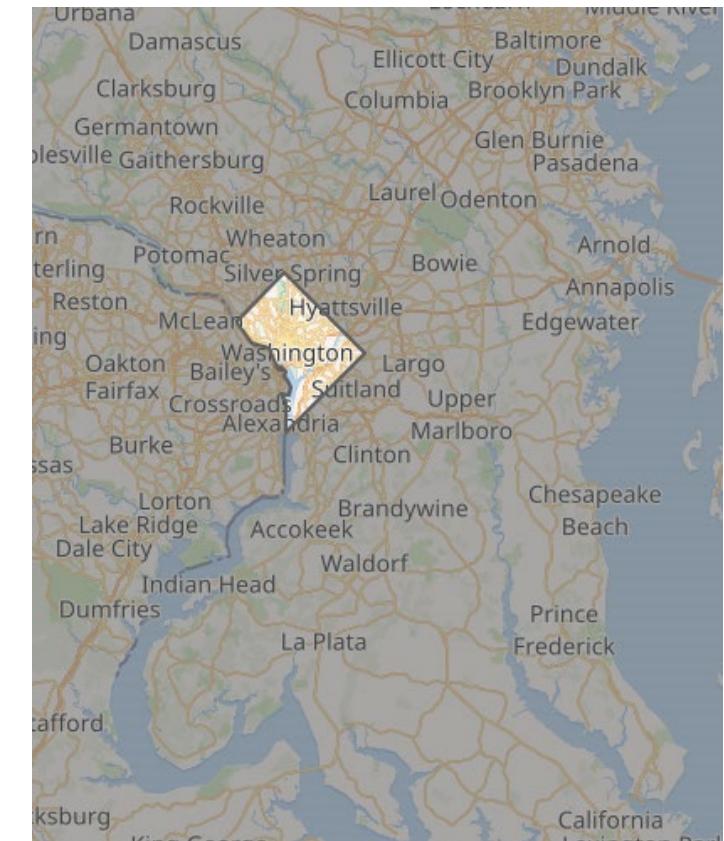
ie., $< 0.00000000000000000000000000000001$. Same conclusion with stronger evidence.

No way that this can happen by chance

Example: Crime data

```
> crime[9,]  
      State murder.rate college  
9 District of Columbia      41.8     38.3
```

- Washington DC, the capital district of USA – is a district, and NOT a state.
- As such, the district does not have the socioeconomic of a state – schools, welfare, etc works differently, for the invited clientele that works there.
- Only 30% of the workforce actually lives there!
- As such, **the data generating mechanism is different!**
- The inclusion on the list is likely for completion, but essentially irrelevant when it comes to the relation between college education and murder rate **at the state level**.



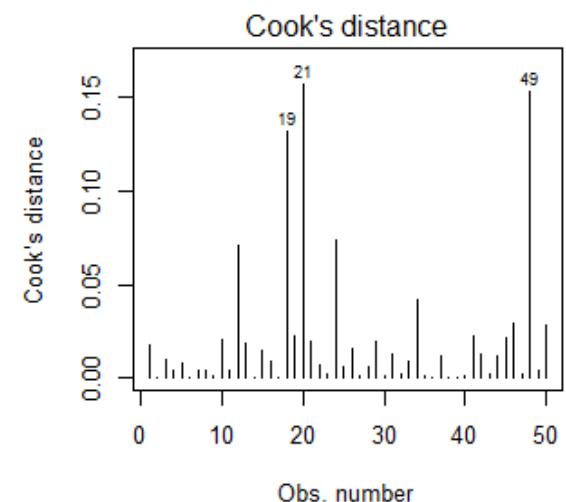
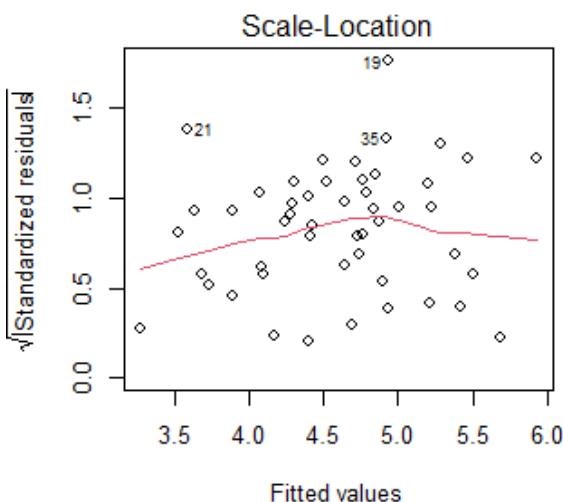
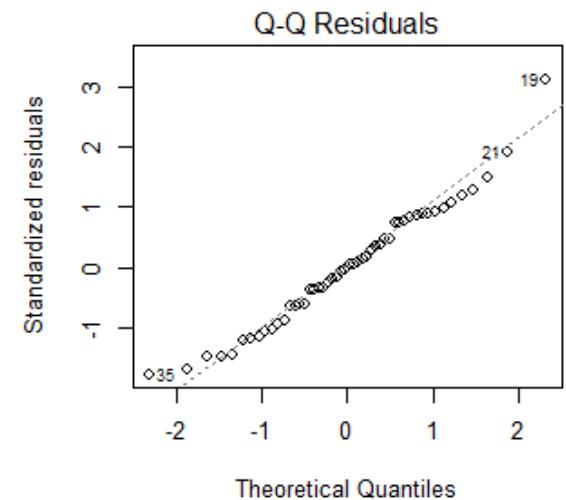
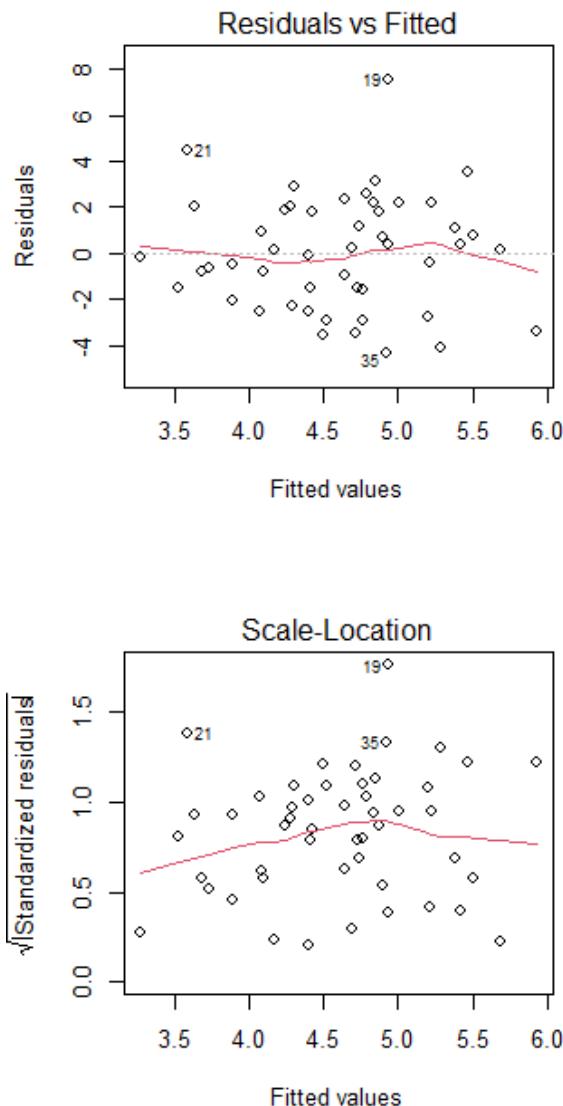
Picture: Wikipedia

We shall stick to STATES, and leave out Washington DC!

Example: Crime data

```
par(mfrow = c(2, 2))
plot(reg2, which = 1:4)
par(mfrow = c(1, 1))
```

- Much nicer; points towards a fitting model: Variance homogeneity and normality.



Example: Crime data

```
> summary(reg2)
```

Call:

```
lm(formula = murder.rate ~ college, data = crime50)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|--------|--------|--------|
| -4.3255 | -1.5533 | 0.0073 | 1.8450 | 7.5607 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|--------------|
| (Intercept) | 8.04159 | 2.06478 | 3.895 | 0.000304 *** |
| college | -0.13788 | 0.08163 | -1.689 | 0.097687 . |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.464 on 48 degrees of freedom

Multiple R-squared: 0.0561, Adjusted R-squared: 0.03644

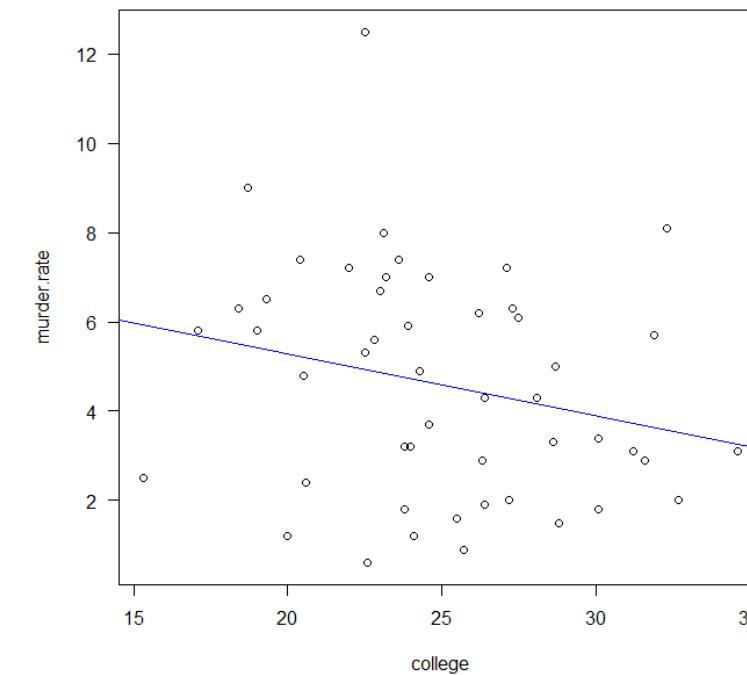
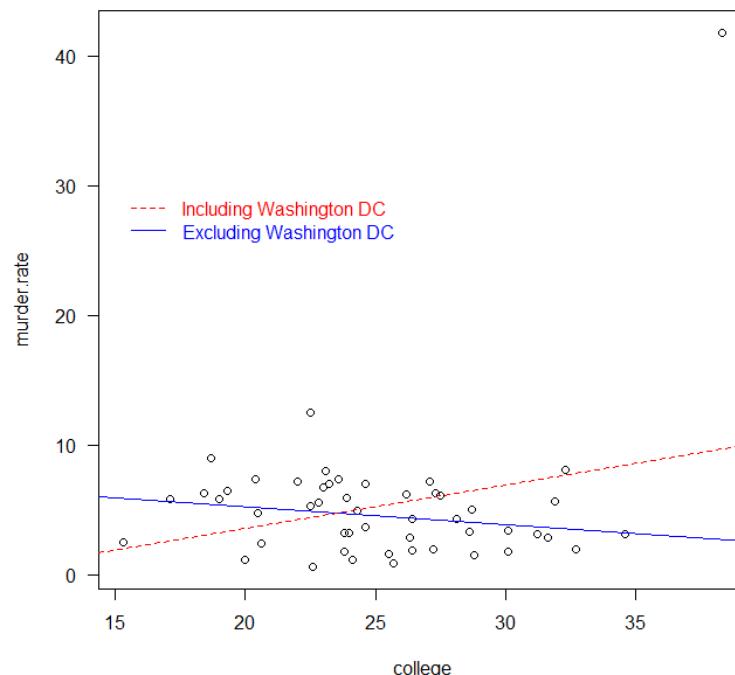
F-statistic: 2.853 on 1 and 48 DF, p-value: 0.09769

Example: Crime data

```
>plot(murder.rate ~ college, data = crime, las=1)
>abline(reg2,font=2,col="blue")
>abline(reg1,font=2,col="red",lty=2)

>legend(15,30,c("Including Washington DC","Excluding Washington DC"),
+        lty=c(2,1),col=c("red","blue"),text.col=c("red","blue"),bty="n")

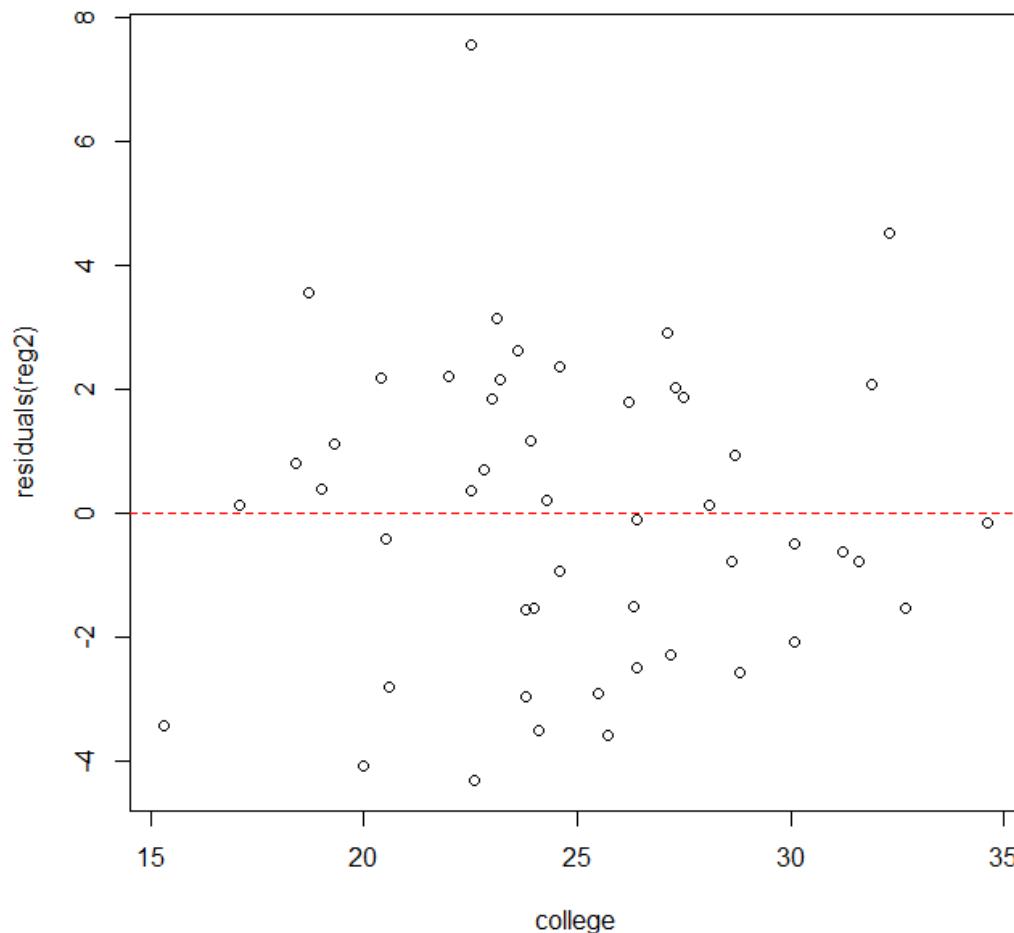
>plot(murder.rate ~ college, data = crime50, las=1)
>abline(reg2,font=2,col="blue")
```



Example: Crime data

- Linearity: Plot residuals vs. numerical explanatory variables:

```
plot(residuals(reg2) ~ college, data=crime50)
abline(h = 0, lty = 2, col="red")
```

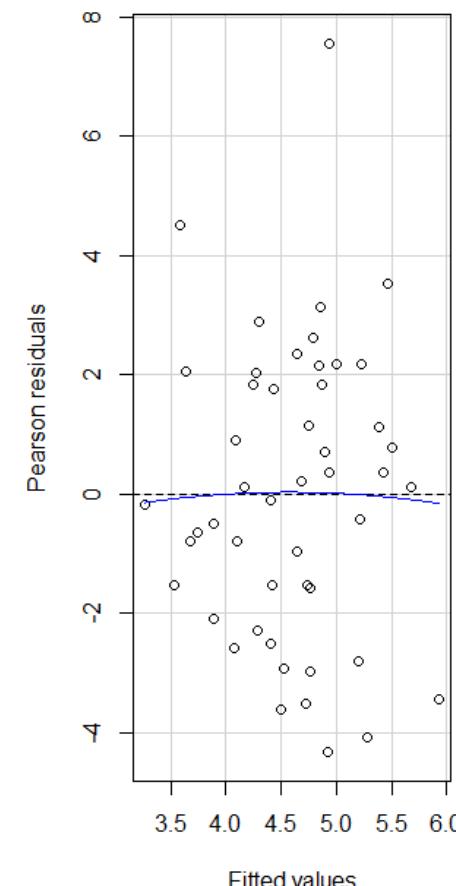
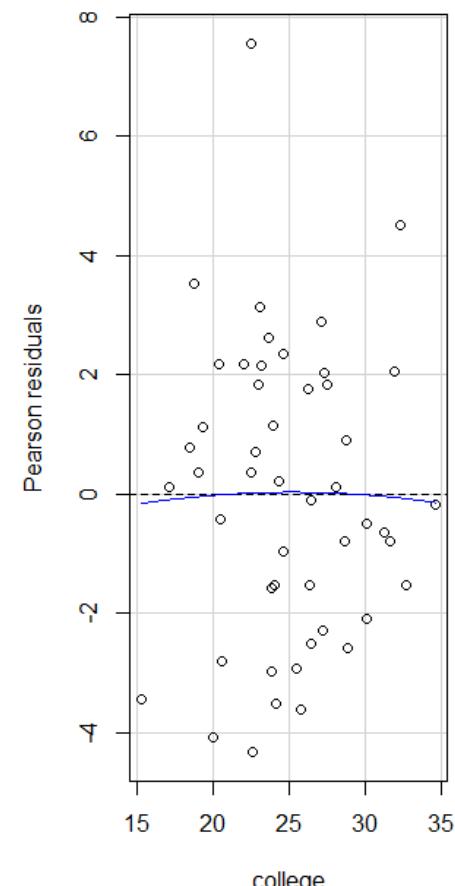


Example: Crime data

- Linearity: Plot residuals vs. numerical explanatory variables with estimated curvature:

```
> residualPlots(reg2)
      Test stat Pr(>|Test stat|)
```

| | | |
|------------|---------|--------|
| college | -0.1277 | 0.8989 |
| Tukey test | -0.1277 | 0.8984 |

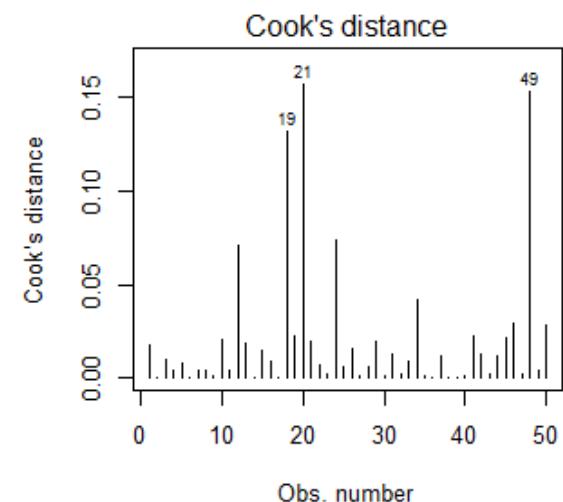
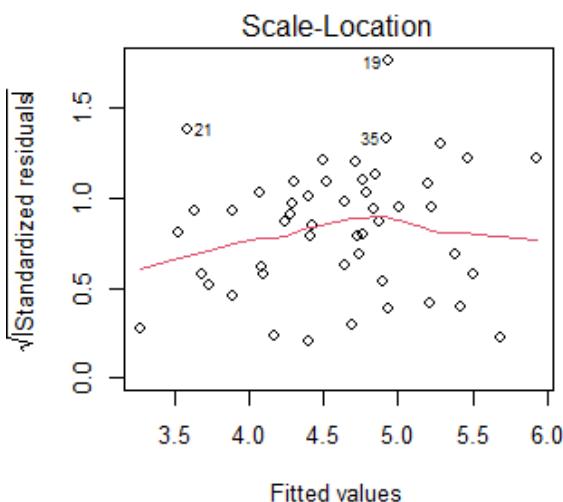
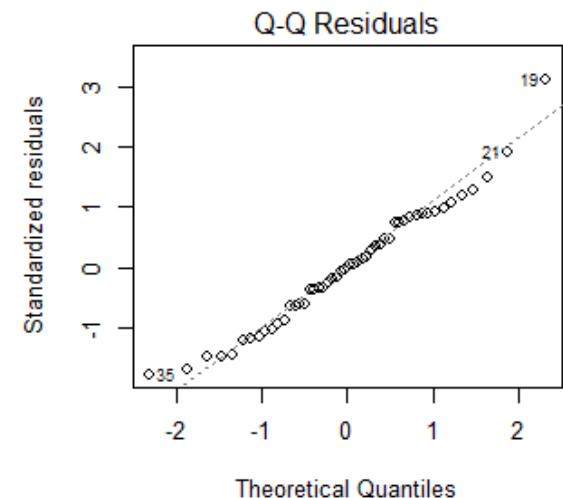
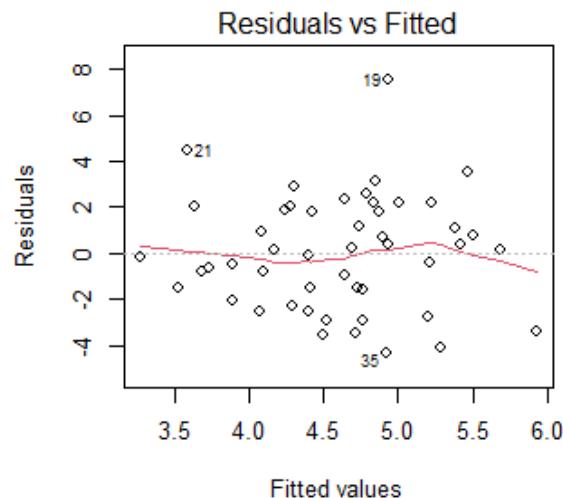


Example: Crime data

```
> # curvature college:  
> reg3<-update(reg2,~.+I(college^2))  
> drop1(reg3,test="F")  
Single term deletions  
  
Model:  
murder.rate ~ college + I(college^2)  
          Df Sum of Sq    RSS    AIC F value Pr(>F)  
<none>            291.34 94.123  
college      1  0.016451 291.35 92.125  0.0027 0.9591  
I(college^2)  1  0.101079 291.44 92.140  0.0163 0.8989  
  
> # curvature fitted values (Tukeys test):  
> reg3<-update(reg2,~.+I(fitted(reg2)^2))  
> drop1(reg3,test="F")  
Single term deletions  
  
Model:  
murder.rate ~ college + I(fitted(reg2)^2)  
          Df Sum of Sq    RSS    AIC F value Pr(>F)  
<none>            291.34 94.123  
college      1  0.42886 291.76 92.196  0.0692 0.7937  
I(fitted(reg2)^2) 1  0.10108 291.44 92.140  0.0163 0.8989
```

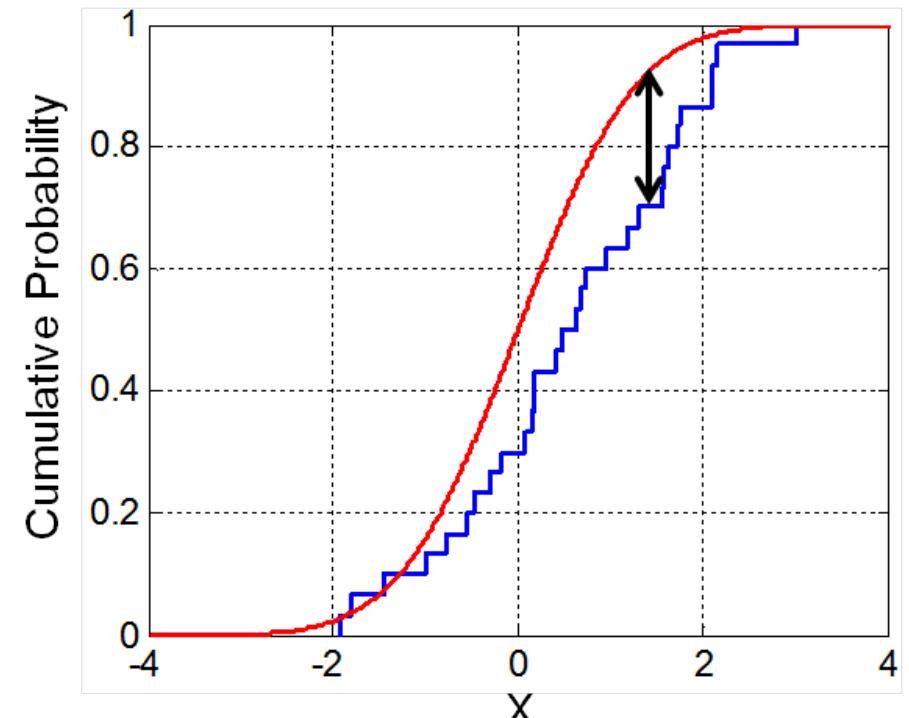
Example: Crime data

- Visual testing – the **EBT** test.
- Why not use one of the many numerical tests for normality?



Example: Crime data

- The *Kolmogorov-Smirnov* test for normality:
- Compares (here) the distance between an empirical and a theoretical distribution function.
- Other tests are the Shapiro-Wilks test and the Anderson-Darling test.
- In **R**: `ks.test()`



Picture source: Wikipedia/Bscan

Tests For Normality

- The Kolmogorov-Smirnov test for normality:

```
> ks.test(rnorm(100), "pnorm")  
  
Asymptotic one-sample Kolmogorov-Smirnov test  
  
data: rnorm(100)  
D = 0.063339, p-value = 0.8172  
alternative hypothesis: two-sided
```

- p=0.82, Test accepts – all good.
- Do normally distributed data follow an exponential distribution?

```
> ks.test(abs(rnorm(100)), "pexp")  
  
Asymptotic one-sample Kolmogorov-Smirnov test  
  
data: abs(rnorm(100))  
D = 0.1113, p-value = 0.1678  
alternative hypothesis: two-sided
```

- p=0.17, Test accepts – not good.

Tests For Normality

- With $X \sim N(0,1)$, $E[|X|] = \sqrt{\frac{2}{\pi}} \approx 0.8$.

What happens if we perturb our data a little? Say, 1% on average:

$$Z = X + 0.008 * Y$$

With Y exponentially distributed (mean 1).

More than 99.99% of the variation comes from X .

```
> ks.test(rnorm(100)+0.008*rexp(100), "pnorm")  
  
Asymptotic one-sample Kolmogorov-Smirnov test  
  
data: rnorm(100) + 0.008 * rexp(100)  
D = 0.094395, p-value = 0.335  
alternative hypothesis: two-sided
```

p=0.34; so little effect that it isn't recognized. All good.

Tests For Normality

- Do the same with many data, more than 100, say 100.000:

```
> ks.test(rnorm(100000)+0.01*rexp(100000), "pnorm")
```

```
Asymptotic one-sample Kolmogorov-Smirnov test
```

```
data: rnorm(1e+05) + 0.008 * rexp(1e+05)
D = 0.0057518, p-value = 0.002676
alternative hypothesis: two-sided
```

- $p=0.003$ significant, even though we on average just perturb with 1% of the average value.

**Even minor, irrelevant model aberrations are detected for large data;
For small data, many incorrect hypotheses are accepted.**

Solution: Graphical tests and common sense.

Example: Crime data

- Post hoc analysis – confidence and prediction intervals:
- The confidence interval for the fitted line at co_0 :

$$\begin{aligned} V(\hat{\alpha} + \hat{\beta}co_0) &= V(\hat{\alpha}) + co_0^2V(\hat{\beta}) + 2co_0Cov(\hat{\alpha}, \hat{\beta}) \\ &= \sigma^2 \left(\frac{1}{n} + \frac{(co_0 - \bar{co})^2}{SSD_{co}} \right), \end{aligned}$$

ie. confidence interval

$$\hat{\alpha} + \hat{\beta}co_0 \pm t_{\frac{\alpha}{2}}(n-2)\hat{\sigma} \sqrt{\frac{1}{n} + \frac{(co_0 - \bar{co})^2}{SSD_{co}}}$$

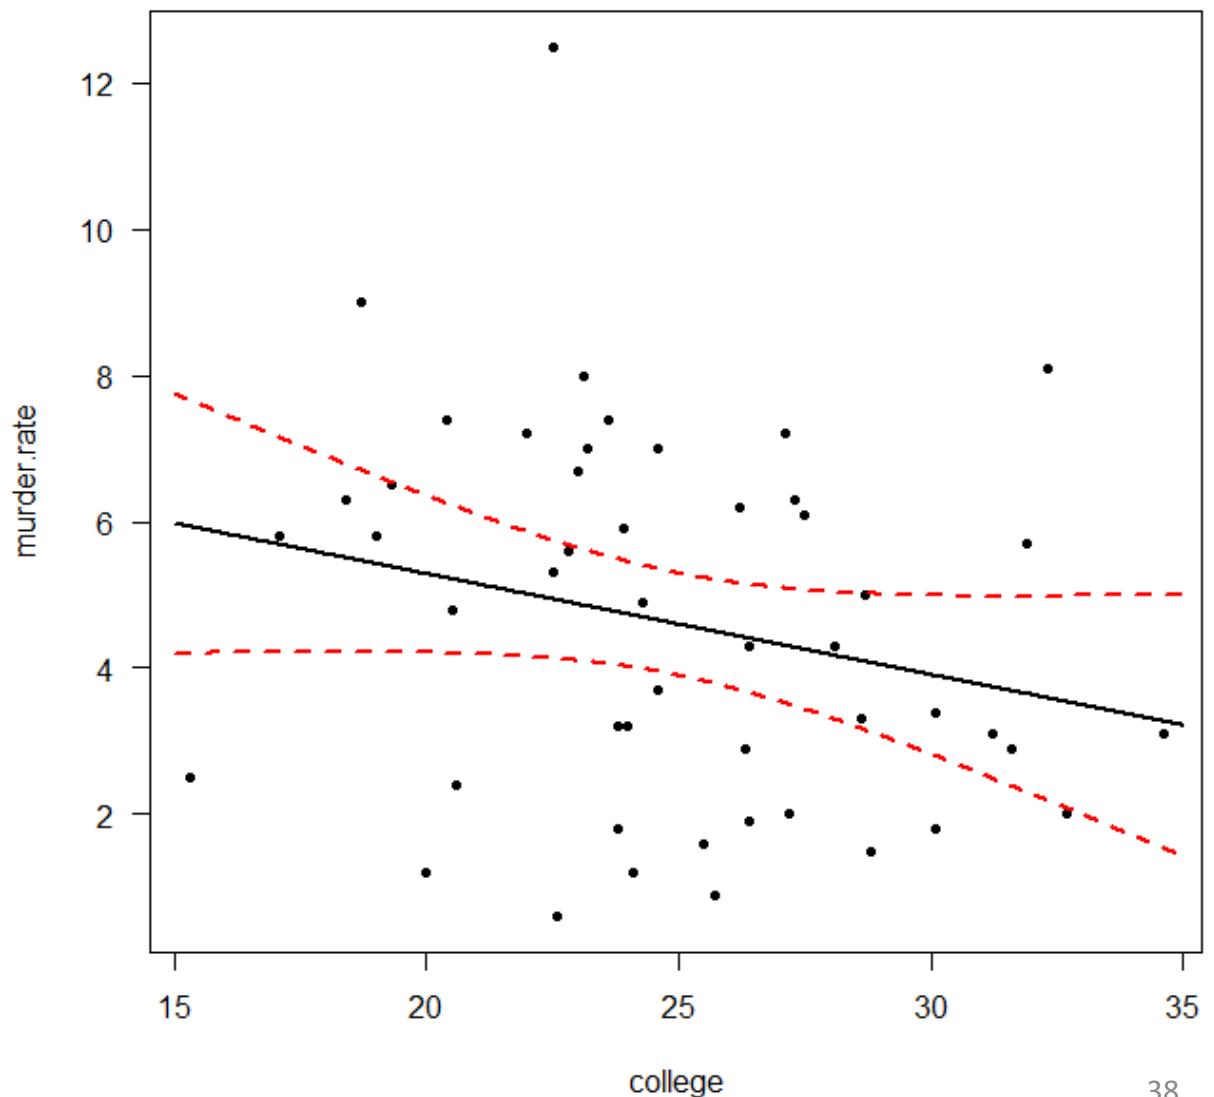
With $t_{\frac{\alpha}{2}}(n-2)$ the $\alpha/2$ percentile in a t-distribution with $n-2$ degrees of freedom.

- Here, $\alpha = 0.05$, $n = 50$, and thus $t_{\frac{\alpha}{2}}(n-2) = -2.010635$.

Example: Crime data

```
>xval <- seq(from=15, to=35, length.out=500)
>newData <- data.frame(college=xval)
>Pred.ci <- predict(reg2, newdata=newData,
+                     interval="confidence",
+                     level=.95)

>plot(murder.rate ~ college, data = crime50, pch = 20, las = 1)
>lines(xval, Pred.ci[, "fit"], lwd=2) ## or use: abline(reg2)
>lines(xval, Pred.ci[, "lwr"], lty=2, col="red", lwd=2)
>lines(xval, Pred.ci[, "upr"], lty=2, col="red", lwd=2)
```



Example: Crime data

- Post hoc analysis – confidence and prediction intervals:
- The prediction interval for a new observation with college rate co_0 :

$$\begin{aligned} V(Y - (\hat{\alpha} + \hat{\beta}co_0)) &= V(Y) + V(\hat{\alpha} + \hat{\beta}co_0) \\ &= \sigma^2 \left(1 + \frac{1}{n} + \frac{(co_0 - \bar{co})^2}{SSD_{co}} \right), \end{aligned}$$

ie. prediction interval

$$\hat{\alpha} + \hat{\beta}co_0 \pm t_{\frac{\alpha}{2}}(n-2)\hat{\sigma} \sqrt{1 + \frac{1}{n} + \frac{(co_0 - \bar{co})^2}{SSD_{co}}}$$

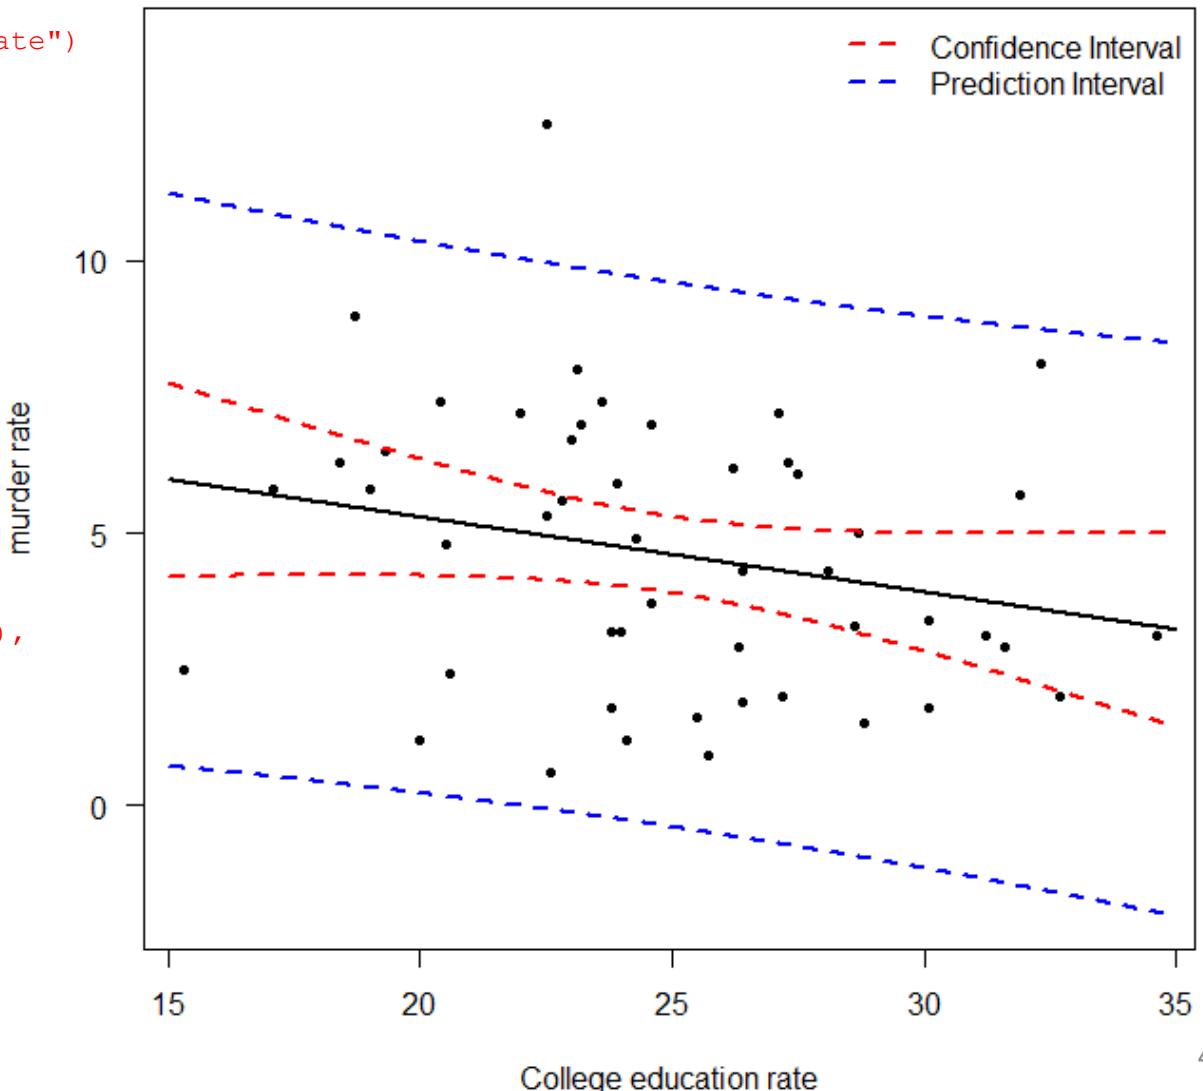
With $t_{\frac{\alpha}{2}}(n-2)$ the $\alpha/2$ percentile in a t-distribution with $n-2$ degrees of freedom.

- Wider than the confidence interval.

Example: Crime data

```
># redoing plot for wider range:  
>plot(murder.rate ~ college, data = crime50, pch = 20, las = 1,  
+      ylim=c(-2,14),xlab="College education rate",ylab="murder rate")  
>lines(xval, Pred.ci[, "fit"], lwd=2)  
>lines(xval, Pred.ci[, "lwr"], lty=2, col="red", lwd=2)  
>lines(xval, Pred.ci[, "upr"], lty=2, col="red", lwd=2)  
  
>## Prediction interval for a new observation:  
>Pred.pi <- predict(reg2, newdata=newData,  
>                     interval="prediction")  
  
>## Add prediction intervals to plot:  
>lines(xval, Pred.pi[, "lwr"], lty=2,  
+       col="blue", lwd=2)  
>lines(xval, Pred.pi[, "upr"], lty=2,  
+       col="blue", lwd=2)  
>legend("topright",c("Confidence Interval","Prediction Interval"),  
+       lty=rep(2,1),lwd=rep(2,2),col=c("red","blue"),bty="n")
```

- Points outside PI: 1
- Expected: 2.5



Example: Crime data

- Influential observations.
- How much does the model change if an observation is left out?
- That is measured with the so-called

$$\bullet \quad D_i = \frac{\|\hat{Y} - \hat{Y}_{(i)}\|^2}{p \hat{\sigma}^2},$$

Where \hat{Y} is the fitted values, $\hat{Y}_{(i)}$ is the fitted values when observation i is left out of the estimation, and p is the rank of the design matrix X .

- With h_{ii} the *leverage* of observation i , the diagonal element of $P_M = X(X^T X)^{-1}X^T$:

$$h_{ii} = X_i(X^T X)^{-1}X_i^T;$$

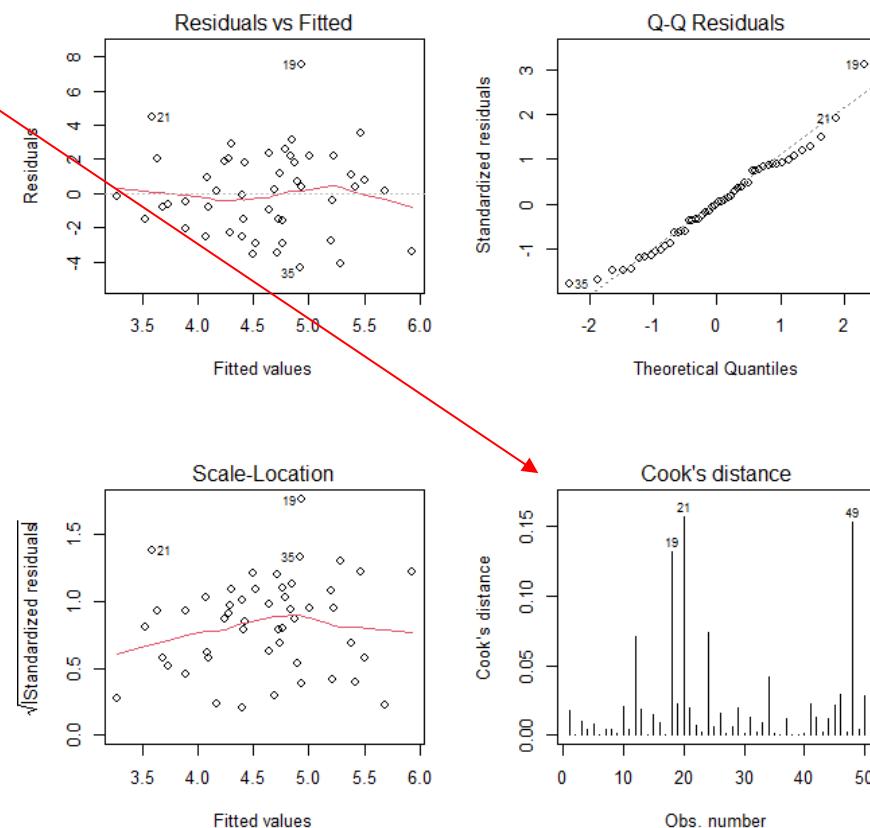
$$e_i = Y_i - \hat{Y}_i,$$

a representation is

$$D_i = \frac{e_i^2}{p \hat{\sigma}^2} \frac{h_{ii}}{(1 - h_{ii})^2}$$

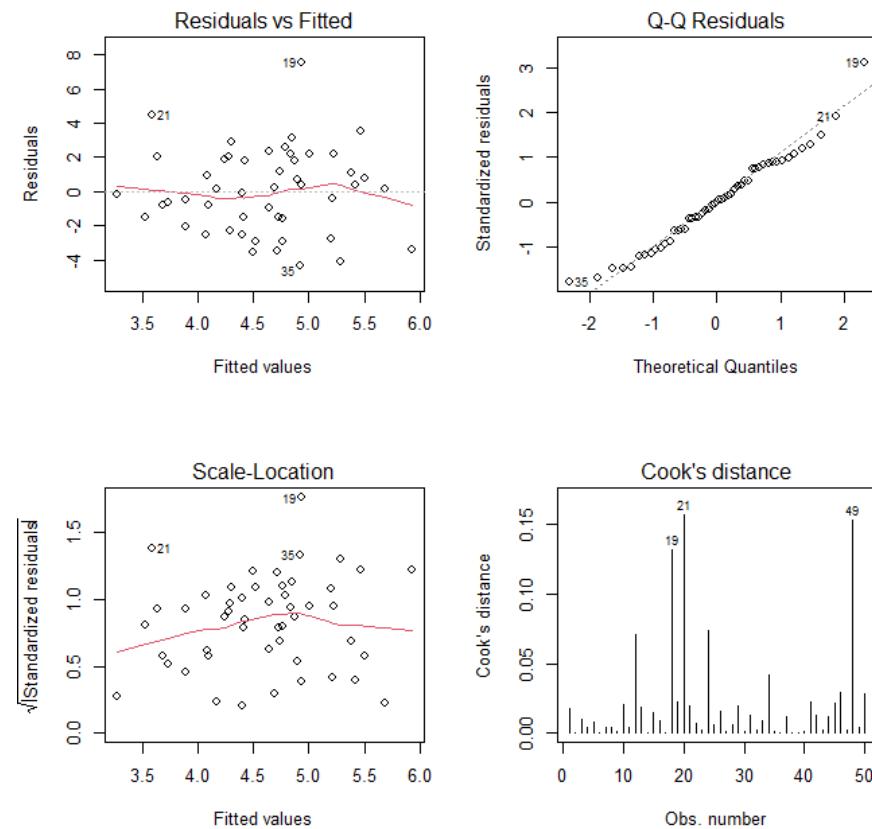
Note that Cook's distance is high when both the leverage h_{ii} and the (squared) residual e_i^2 is high.

Cook's distance D_i :



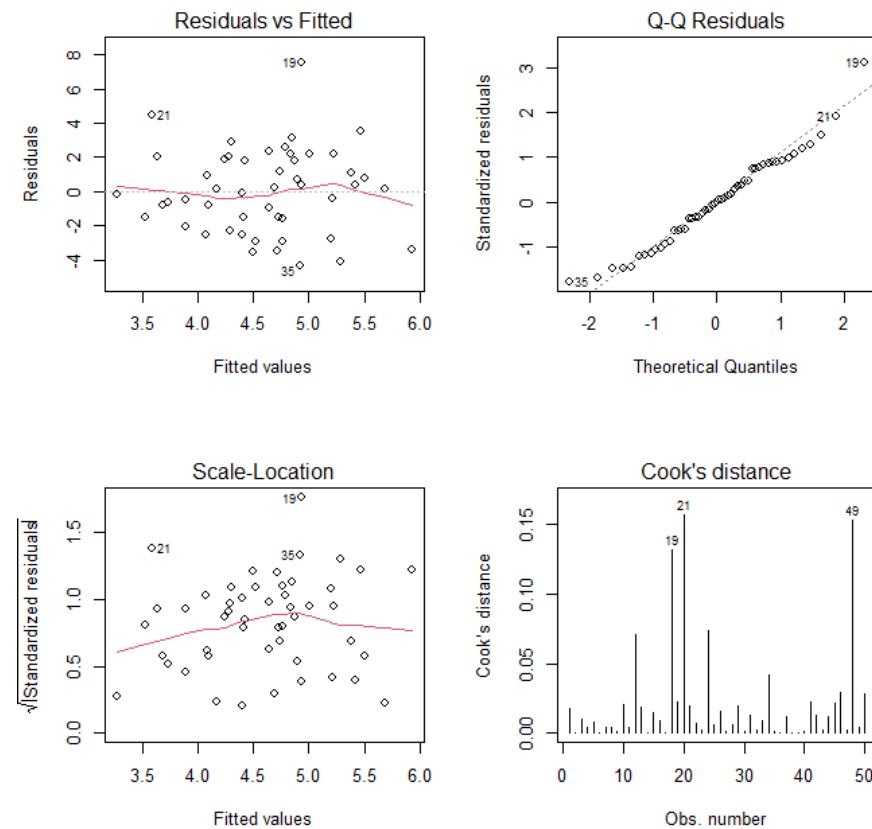
Example: Crime data

- When Cook's distance is high for an observation, results are affected by that single observation.
- If conclusions rely on the presence of single observations, the conclusions are *fragile*.
- We call an observation with a high Cook's distance for an *influential observation*.
- Influential observations should be investigated to see if they are indeed correct, and due to insight into the data generating mechanism; but they are valid data if they are not outliers, and important for the conclusion.
- In the presence of influential observations, reservations in the conclusion need to be made.



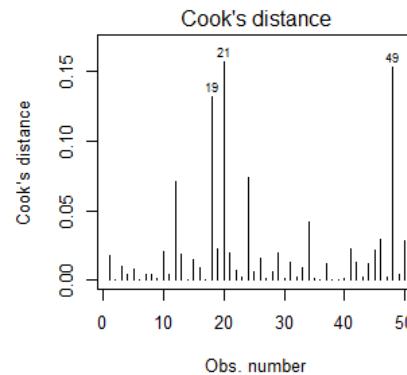
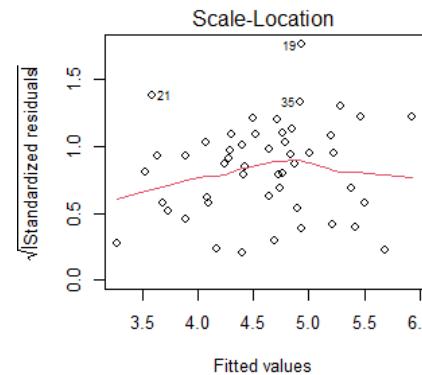
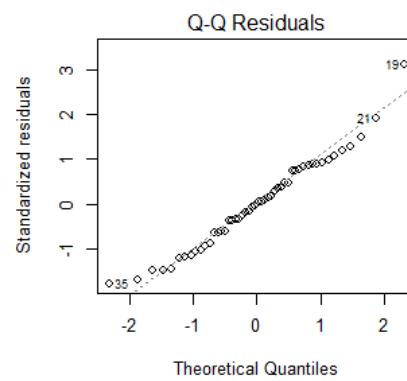
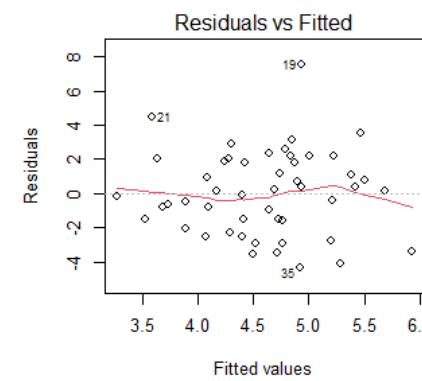
Example: Crime data

- Outliers are often influential observations.
- Outliers are **ALWAYS** corrected or removed;
- Influential observations which are not outliers are subject to investigation and possible correction, but they are **NEVER** removed.
- Reason: No indication of a different data generating mechanism.



Example: Crime data

- When is Cook's distance big?
- No universal agreement about that.
- **Advise:** Do **not** spend time on size dependent formulas like $D > \frac{1}{n}$, $D > \frac{1}{n-p}$ etc.
- If $D > 1$ you should *always* investigate; if $D > \frac{1}{2}$ you should *likely* investigate; If $D < \frac{1}{2}$ *don't bother*.
- With the crime data, a $\frac{1}{n}$ or $\frac{1}{n-p}$ cut-off would leave 5 states for investigation. A $\frac{1}{2}$ cut-off leaves none. We will investigate none.



Example: Crime data

| | Estimate | Std..Error | Lower | Upper | p.value |
|-------------|----------|------------|-------|-------|---------|
| (Intercept) | 8.04 | 2.06 | 3.89 | 12.19 | <0.001 |
| college | -0.14 | 0.08 | -0.30 | 0.03 | 0.0977 |

The fitted model relating murder rate in US states with college education rates has the form:

$$murder.rate_i = 8.04 - 0.14college_i + \varepsilon_i, \varepsilon_i \sim N(0, 2.46^2), i = 1, \dots, n$$

- Washington DC was left out of the data as it is not a state, and it exhibited a different data generating mechanism.

Example: Timber Hardness

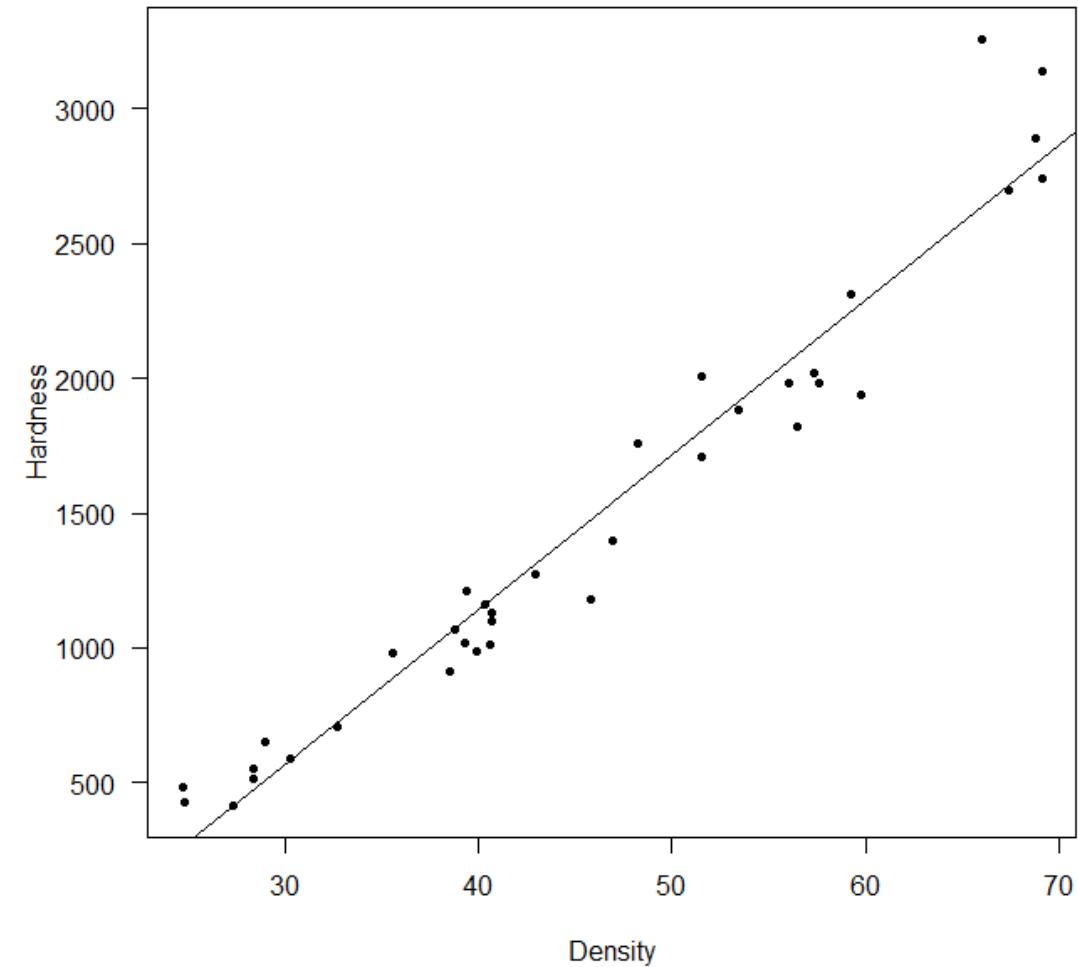
- 36 observations of timber hardness and density from trees in Australia.
- The aim of the study: To estimate parameters that determine the relationship between timber hardness and density, to predict (unknown) timber hardness in future samples based on the simpler measurement of timber density.
- First model: Simple regression.

$$Hardness_i = \alpha + \beta Density_i + \varepsilon_i, \varepsilon_i \text{ iid } \sim N(0, \sigma^2), i = 1, \dots, 36.$$

- Does this model fit?

Example: Timber Hardness

```
>janka <- read.table("Data/janka.txt", header=TRUE, quote="\"")  
> names(janka) <- c("Density", "Hardness")  
> summary(janka)  
    Density      Hardness  
Min.   :24.70   Min.   : 413.0  
1st Qu.:37.77   1st Qu.: 962.8  
Median :41.80   Median :1195.0  
Mean   :45.73   Mean   :1469.5  
3rd Qu.:56.70   3rd Qu.:1980.0  
Max.   :69.10   Max.   :3260.0  
  
> reg3 <- lm(Hardness ~ Density, data = janka)  
> plot(Hardness ~ Density, data = janka, pch = 20, las = 1)  
> abline(reg3)
```

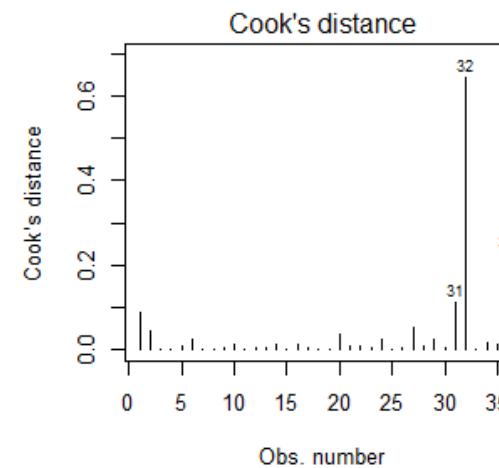
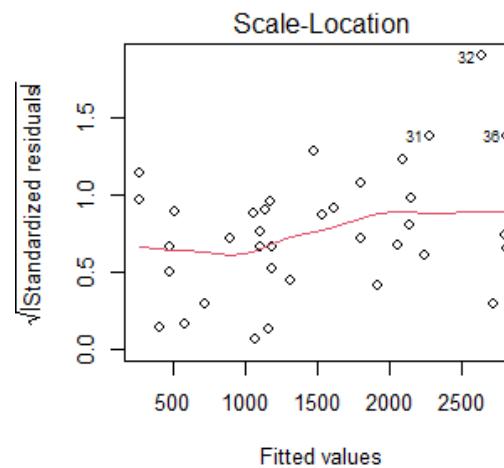
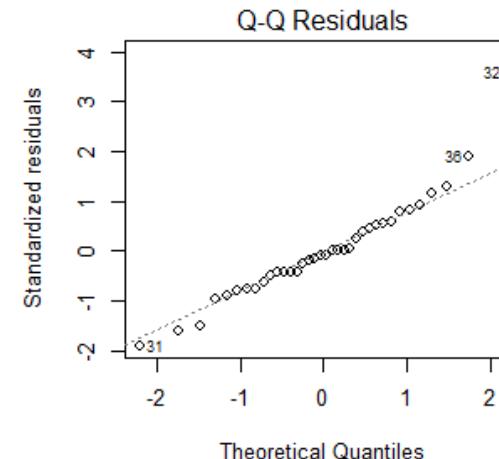
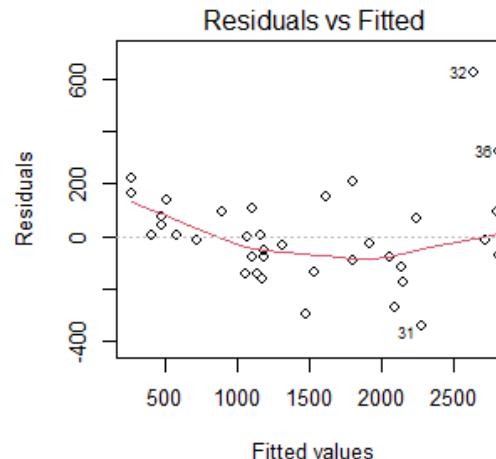
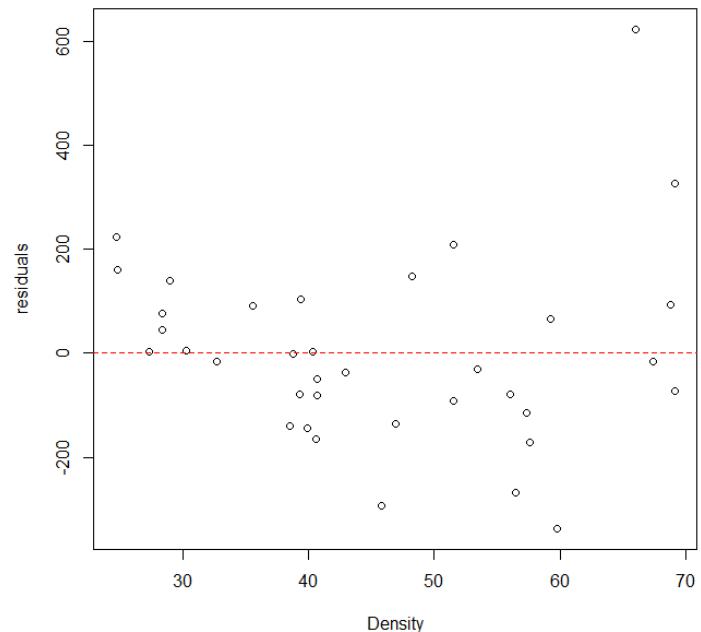


Example: Timber Hardness

- Model fit:

```
> par(mfrow=c(2, 2))
> plot(reg3, which=1:4)
> par(mfrow=c(1, 1))

> plot(residuals(reg3) ~ Density, data = janka,
+       ylab="residuals")
> abline(h = 0,col="red",lty=2)
```



Example: Timber Hardness

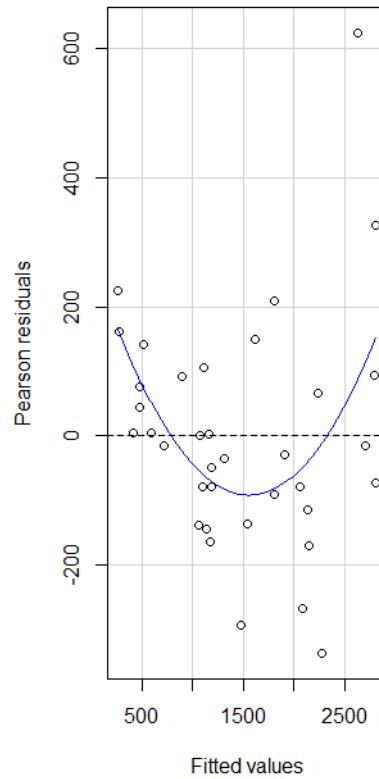
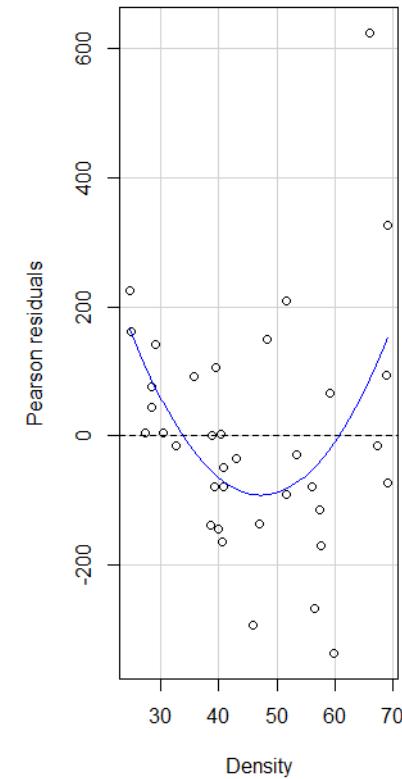
- Model fit:

```
> residualPlots(reg3)

    Test stat Pr(>|Test stat|)

Density      3.2483      0.002669 ** 
Tukey test   3.2483      0.001161 ** 

---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 0.1 ' ' 1
```



Example: Timber Hardness

- The lognormal distribution:

$$X = e^Z, Z \sim N(\mu, \sigma^2)$$

Coefficient of variation:

$$CV = \frac{\sqrt{V(X)}}{E(X)} = \frac{\sqrt{(e^{\sigma^2} - 1)e^{2\mu+\sigma^2}}}{e^{\mu+\frac{1}{2}\sigma^2}} = \sqrt{(e^{\sigma^2} - 1)}$$

Independent of μ when σ^2 is fixed.

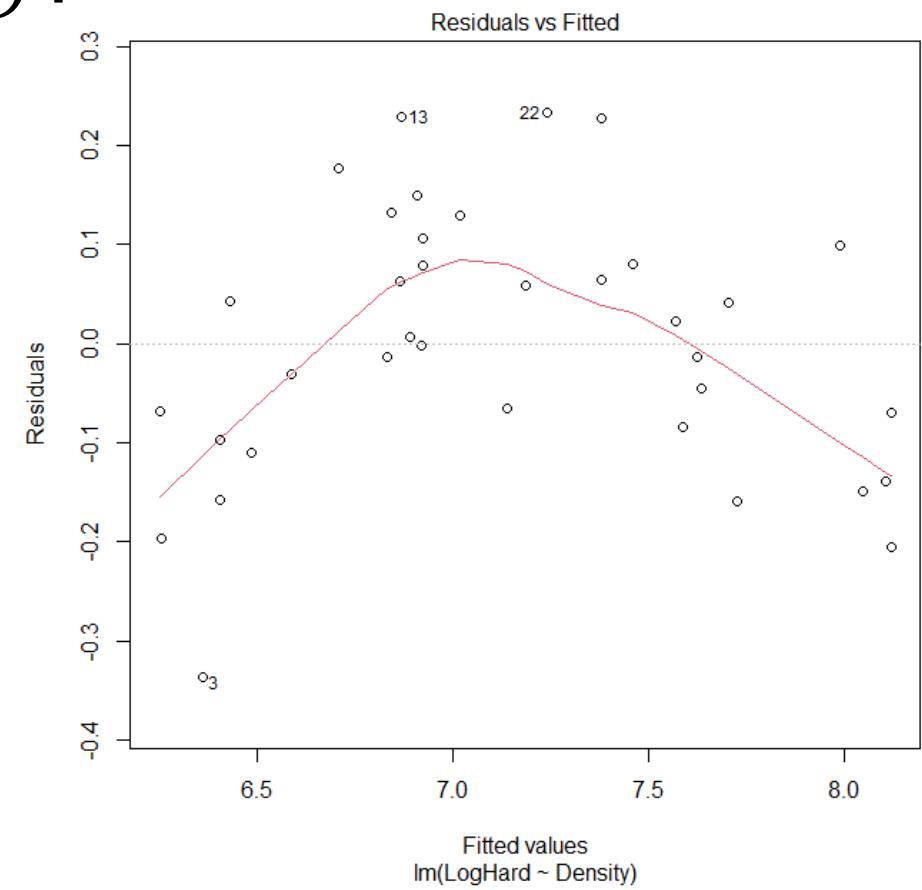
- For fixed σ^2 , SD varies proportional with the mean!

Example: Timber Hardness

- Data transformation: Model $\log(Y)$:

```
janka$LogHard <- log(janka$Hardness)  
reg4 <- lm(LogHard ~ Density, data = janka)  
plot(reg4, which=1)
```

- No longer a trumpet shape, but still curvature.



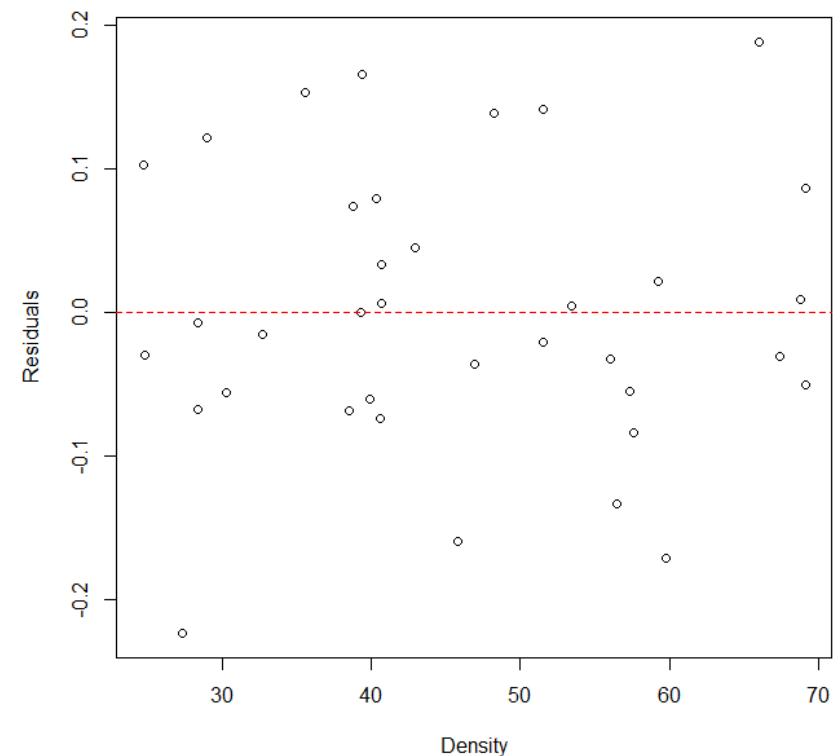
Example: Timber Hardness

- 2nd order Taylor approximation:

$$\log(Y_i) = \alpha + \beta_1 \text{density}_i + \beta_2 \text{density}_i^2 + \varepsilon_i, i = 1, \dots, n$$

```
> reg5 <- lm(LogHard ~ Density + I(Density^2), data = janka)
> plot(residuals(reg5) ~ Density, data = janka,
+       ylab="Residuals")
> abline(h = 0,lty=2,col="red")
```

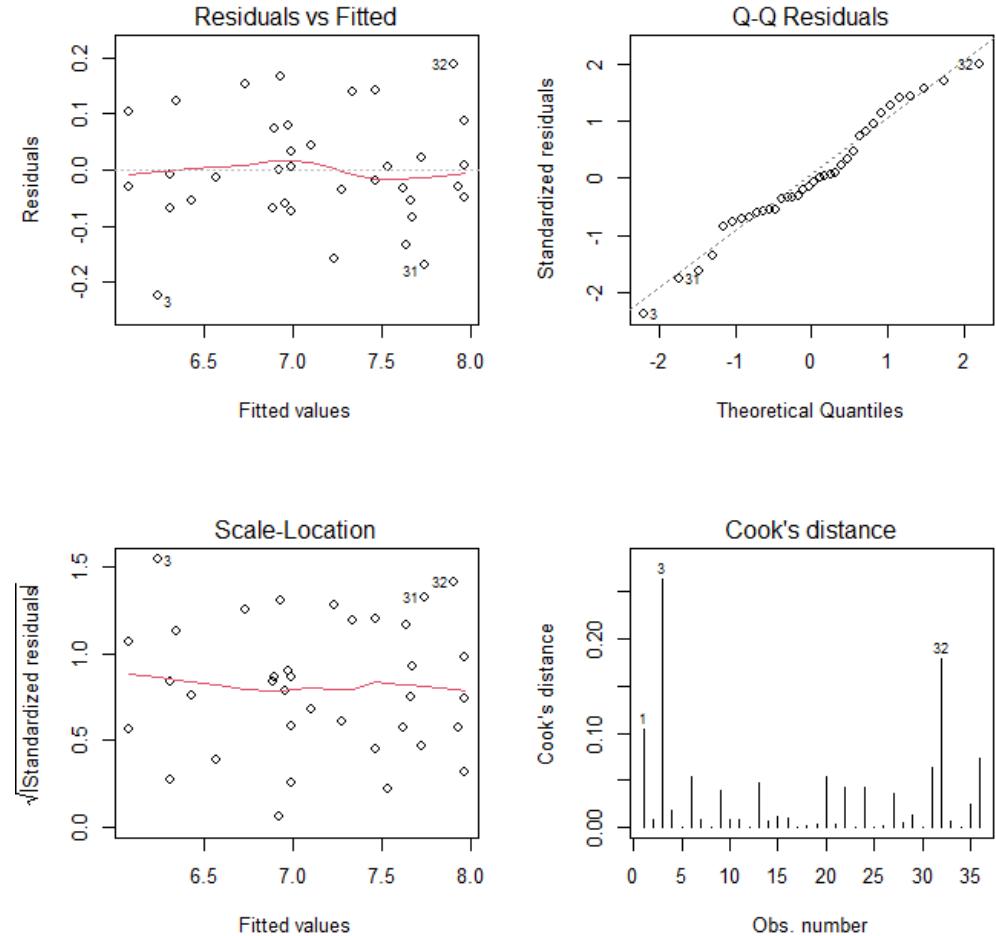
- Much better...



Example: Timber Hardness

```
par(mfrow=c(2, 2))
plot(reg5, which=1:4)
par(mfrow=c(1, 1))
```

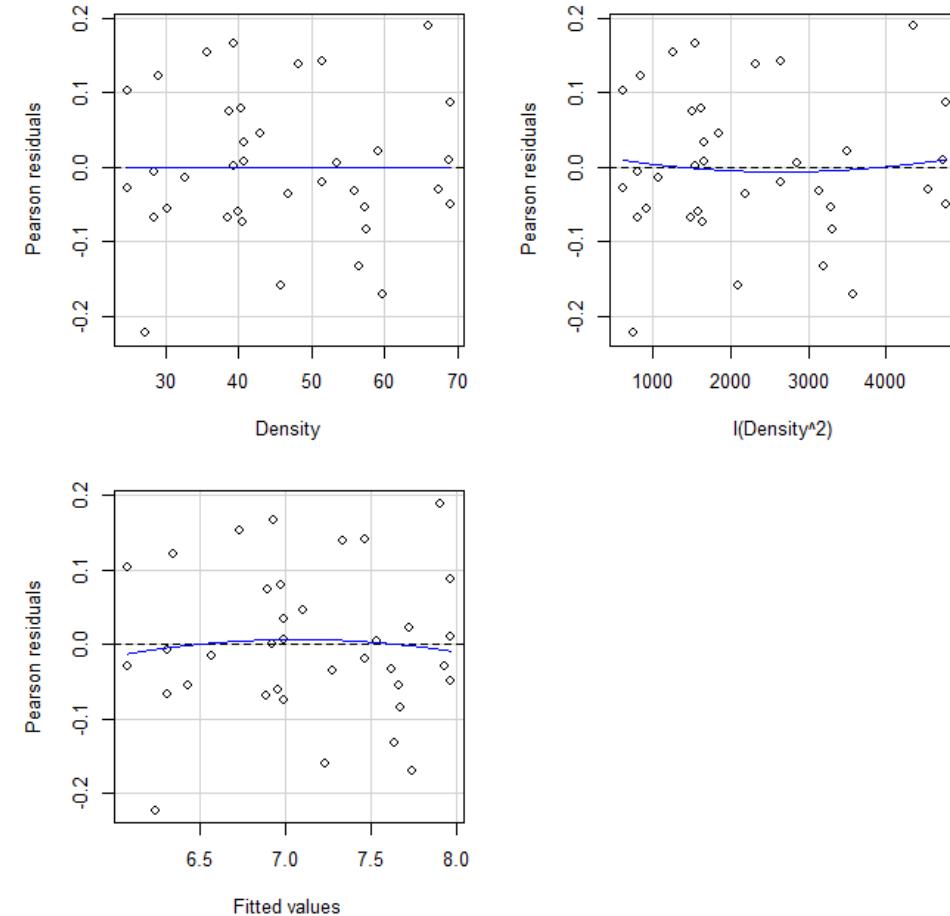
- Variance homogeneity, linearity
(previous slide), normality and no
apparent influential observations.



Example: Timber Hardness

```
> residualPlots(reg5)
      Test stat Pr(>|Test stat|)
Density      -2.4063    0.02206 *
I(Density^2) 1.3611    0.18300
Tukey test   -1.2757    0.20207
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Looks okay, confirming the need for the square term.



Example: Timber Hardness

```
> summary(reg5)

Call:
lm(formula = LogHard ~ Density + I(Density^2), data = janka)

Residuals:
    Min      1Q  Median      3Q     Max 
-0.22331 -0.05708 -0.01104  0.07500  0.18871 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 4.138e+00 2.087e-01 19.828 < 2e-16 ***
Density     9.152e-02 9.305e-03  9.835 2.45e-11 ***
I(Density^2) -5.228e-04 9.764e-05 -5.354 6.49e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1008 on 33 degrees of freedom
Multiple R-squared:  0.9723, Adjusted R-squared:  0.9706 
F-statistic: 578.9 on 2 and 33 DF,  p-value: < 2.2e-16

> drop1(reg5,test="F")
Single term deletions

Model:
LogHard ~ Density + I(Density^2)
          Df Sum of Sq   RSS   AIC F value    Pr(>F)    
<none>           0.33510 -162.37
Density       1   0.98221 1.31731 -115.08  96.725 2.452e-11 ***
I(Density^2)  1   0.29111 0.62621 -141.86  28.668 6.486e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Example: Timber Hardness

Final model:

$$\log(Y_i) = \alpha + \beta_1 \text{density}_i + \beta_2 \text{density}_i^2 + \varepsilon_i, i = 1, \dots, n$$

On the original scale:

$$Y_i = a \cdot b_1^{\text{density}_i} \cdot b_2^{\text{density}_i^2} \cdot \tilde{\varepsilon}_i, i = 1, \dots, n$$

with

$$a = e^\alpha, b_1 = e^{\beta_1}, b_2 = e^{\beta_2}, \log(\tilde{\varepsilon}_i) \text{ iid, } \sim N(0, \sigma^2).$$

- Multiplicative model!

Example: Timber Hardness

Final model:

$$\log(Y_i) = 4.14 + 0.0915 \text{density}_i - 0.000523 \text{density}_i^2 + \varepsilon_i, i = 1, \dots, n$$

On the original scale:

$$Y_i = 62.80 \cdot 1.0958^{\text{density}_i} \cdot 0.9994771^{\text{density}_i^2} \cdot \tilde{\varepsilon}_i, i = 1, \dots, n$$

With

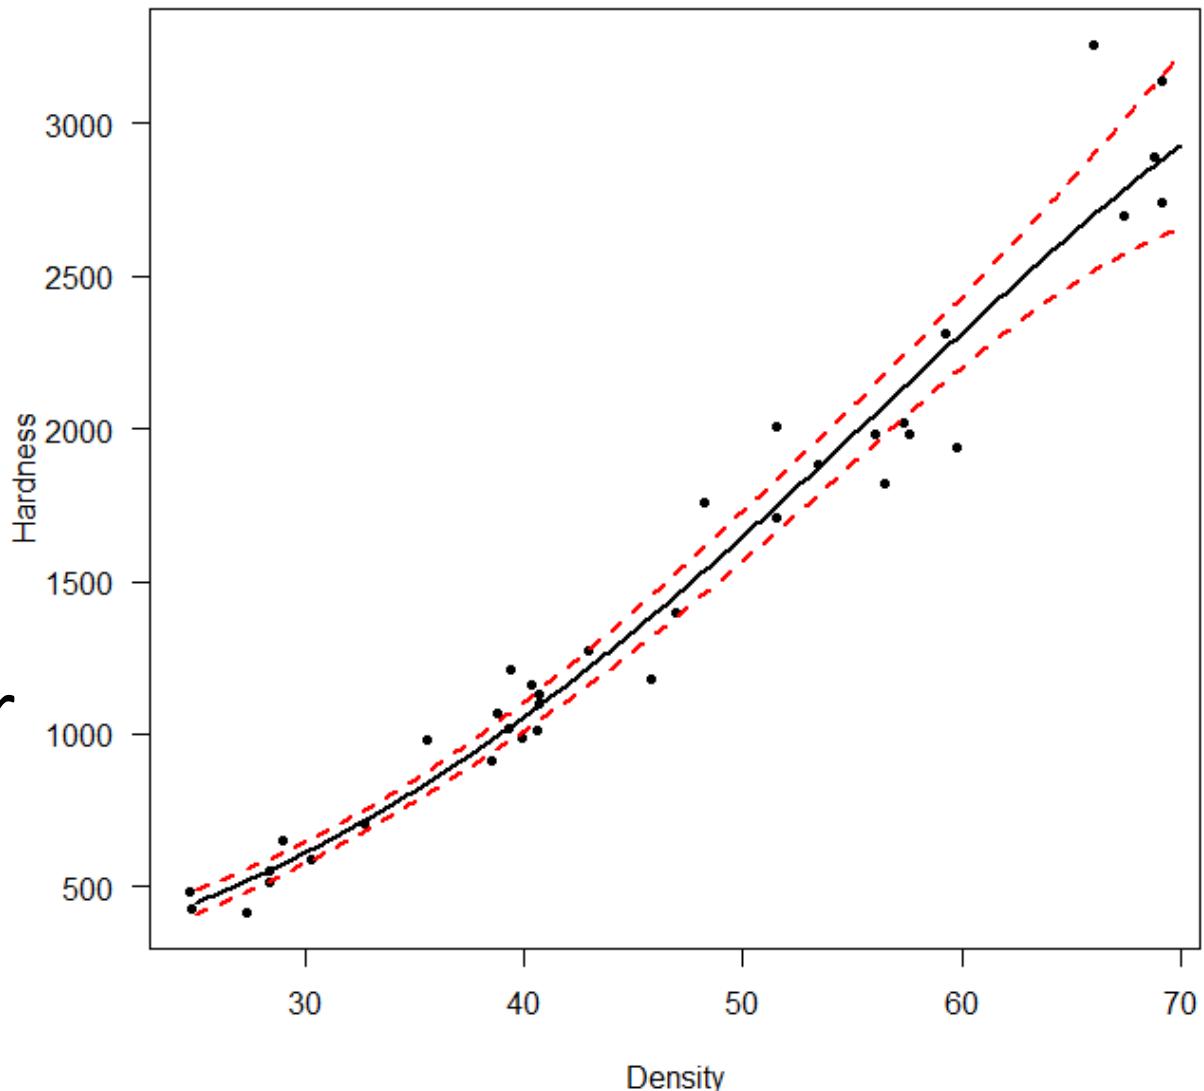
$$62.80 = e^{4.14}, 1.0958 = e^{0.0915}, 0.9994 = e^{-0.000523}, \log(\tilde{\varepsilon}_i) \text{ iid, } \sim N(0, 0.1008^2).$$

Example: Timber Hardness

```
>xval <- seq(from=25, to=70, length.out=500)
>newData <- data.frame(Density=xval)
>Pred.ci <- predict(reg5, newdata=newData,
+                      interval="confidence",
+                      level=.95)

>plot(Hardness ~ Density, data = janka, pch = 20, las = 1)
>lines(xval, exp(Pred.ci[, "fit"]), lwd=2)
>lines(xval, exp(Pred.ci[, "lwr"]), lty=2, col="red", lwd=2)
>lines(xval, exp(Pred.ci[, "upr"]), lty=2, col="red", lwd=2)
```

- Confidence interval increases for large density values but not for small.



Example: Timber Hardness

- Note that

$$E(\log N(\mu, \sigma^2)) = e^{\mu + \sigma^2/2}$$

Strictly, we should have

$$\hat{Y} = \exp(\hat{Z} + V(\hat{Z})/2)$$

and not just

$$\hat{Y} = \exp(\hat{Z})$$

as on the previous slide.

- But the difference is often immaterial, and indeed it is so here:

```
> my.se.fit<-predict(reg5, newdata=newData,
+                      interval="confidence",
+                      level=.95, se.fit=T)$se.fit

> summary(exp(my.se.fit^2/2))
   Min. 1st Qu. Median      Mean 3rd Qu.      Max.
1.000   1.000   1.000   1.000   1.000   1.001
```

- Multiplication factor: 1.

Exercises

- exam 2009, problem 3 (testing).
- Exam 2012 problems 5.1-5.4 (estimation).
- Brain Exercise, including extra questions (model selection).
- Process Exercise (first look at interactions).